

Artificial Intelligence (CS303)

Lecture 3: Adversarial Search

Hints for this lecture

- The *environment* may, in many cases, involves more than one agents.
- *Adversarial* means agents compete to maximize their own benefits.

Outline of this lecture

- Games
- Alpha-Beta Pruning
- Generalization of Minimax Algorithm

I. Games

What is Game?

- More than 1 agent, with conflicting goals.
- We consider two players, zero-sum games.

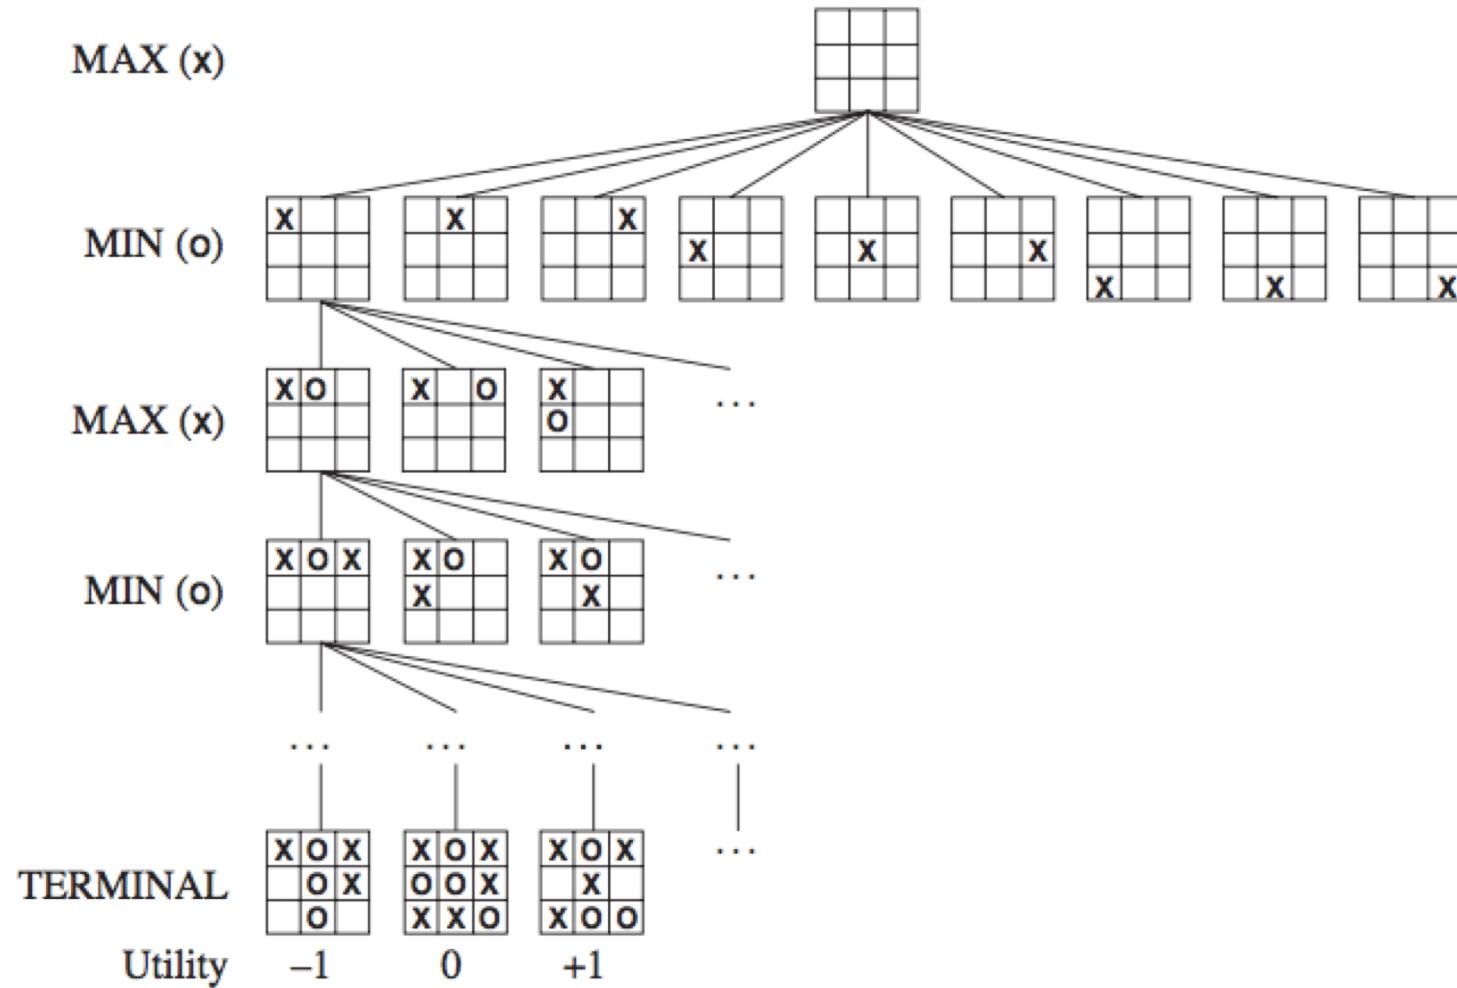


Definition of a game

- S_0 : The **initial state**, which specifies how the game is set up at the start.
- $\text{PLAYER}(s)$: Defines which player has the move in a state.
- $\text{ACTIONS}(s)$: Returns the set of legal moves in a state.
- $\text{RESULT}(s, a)$: The **transition model**, which defines the result of a move.
- $\text{TERMINAL-TEST}(s)$: A **terminal test**, which is true when the game is over and false otherwise. States where the game has ended are called **terminal states**.
- $\text{UTILITY}(s, p)$: A **utility function** (also called an objective function or payoff function),

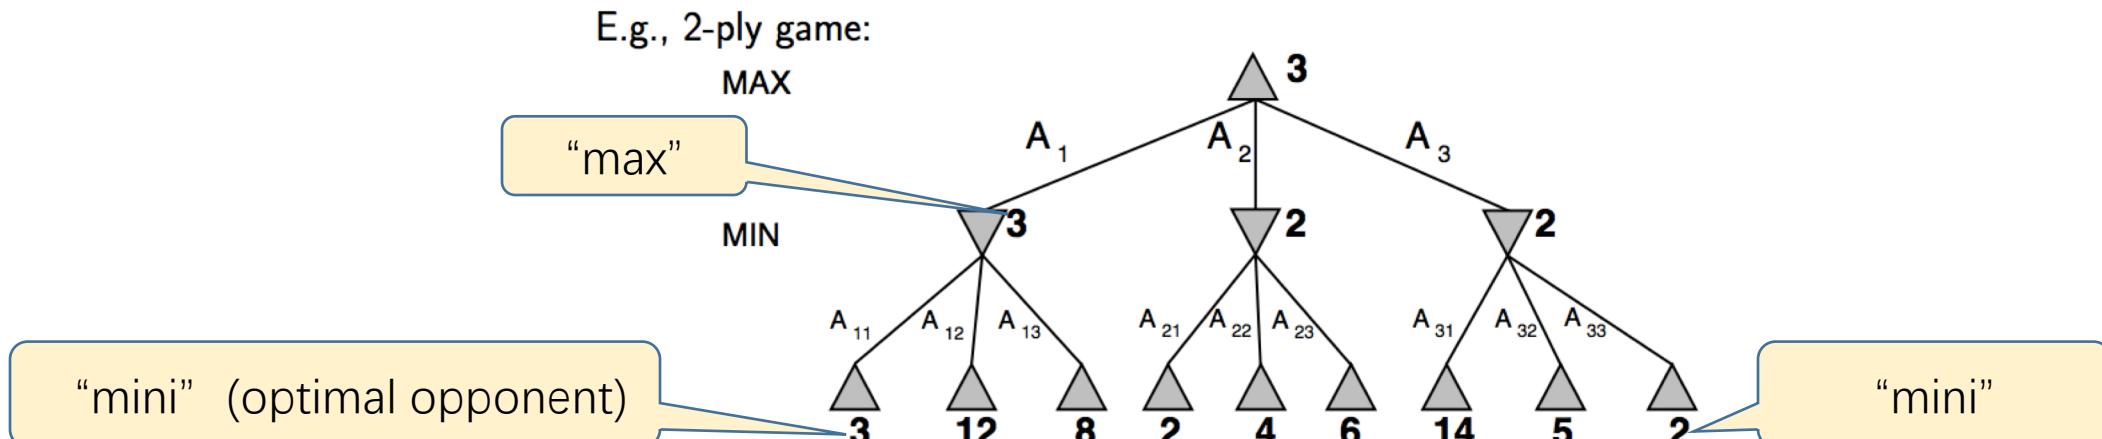
Two players: MAX and MIN

Tic-Tac-Toe Search Tree



Optimal decision in games

- Assume the game is deterministic and perfect information is available
- Idea (for MAX): choose the move to position **the highest minimax value**



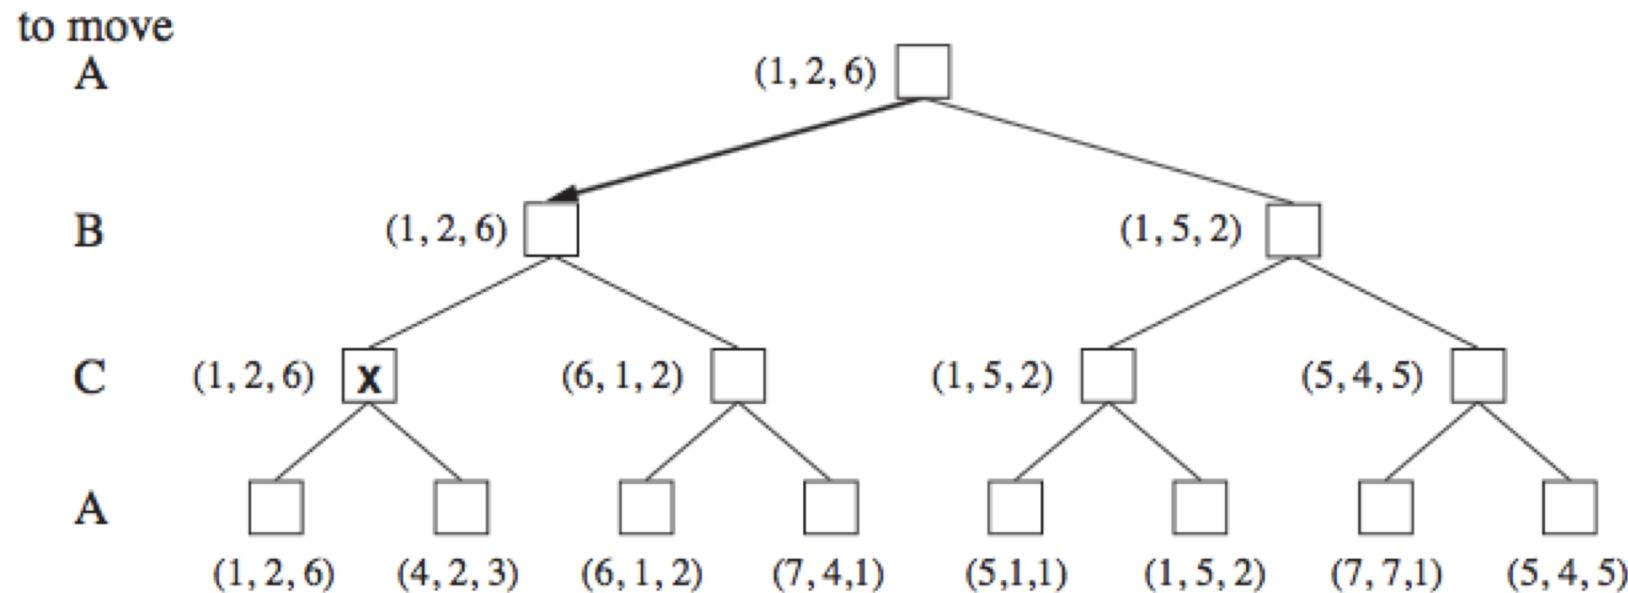
$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Minimax Algorithm

- Perform a complete depth-first search of the game tree.
- Recursively compute the minimax values of each successor state.
- Maximize the worst-case outcome for MAX.

Minimax Algorithm

- Use a vector to represent the utility in case of more than 2 players.

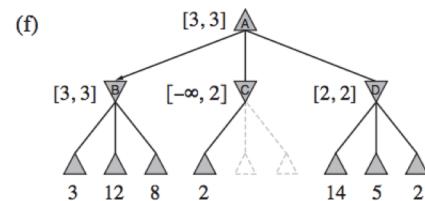
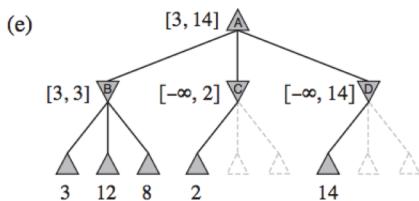
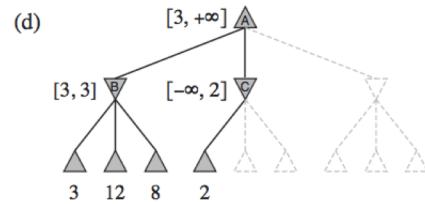
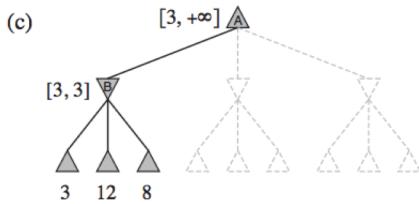
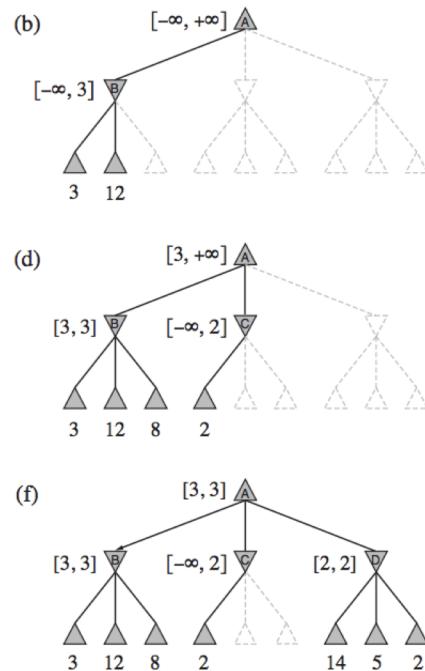
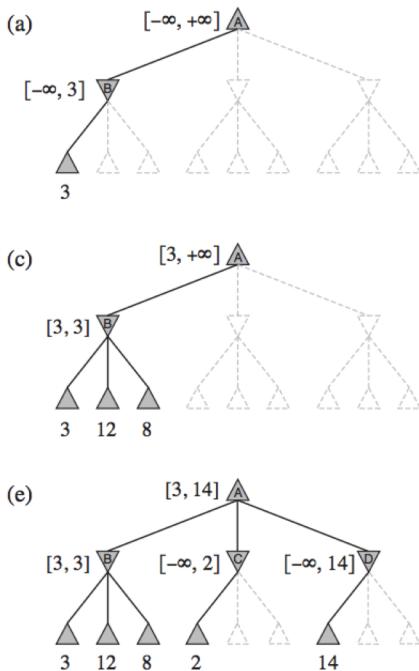


Is minimax algorithm still optimal in this case?

II. Alpha-Beta Pruning

Alpha-Beta Pruning

- A simplification of minimax algorithm.
- Remove (unneeded) part of the minimax tree from consideration.



Alpha: the value of the best (i.e., **highest-value**) choice we have found so far at any choice point along the path for **MAX**.

Beta: the value of the best (i.e., **lowest-value**) choice we have found so far at any choice point along the path for **MIN**.

Alpha-Beta Pruning

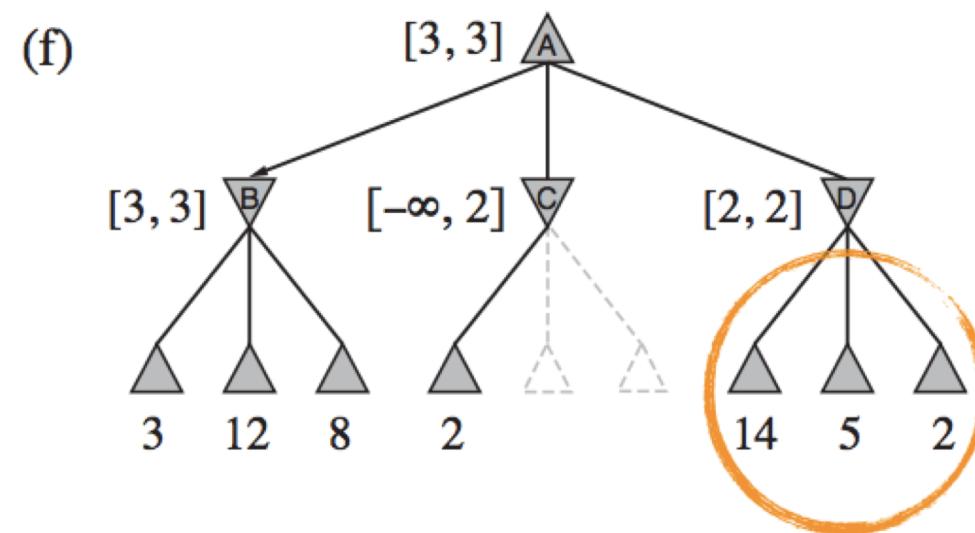
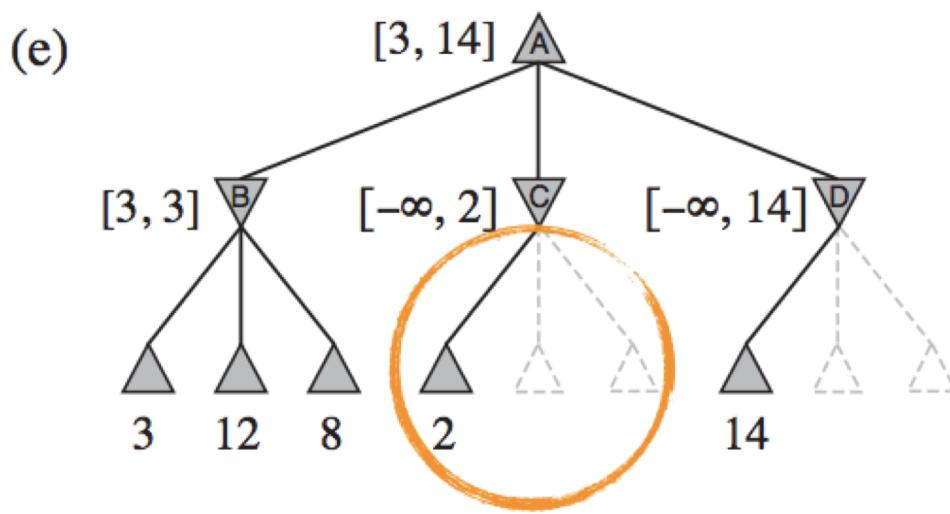
```
function ALPHA-BETA-SEARCH(state) returns an action
    v  $\leftarrow$  MAX-VALUE(state,  $-\infty$ ,  $+\infty$ )
    return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
    if TERMINAL-TEST(state) then return UTILITY(state)
    v  $\leftarrow$   $-\infty$ 
    for each a in ACTIONS(state) do
        v  $\leftarrow$  MAX(v, MIN-VALUE(RESULT(s,a),  $\alpha$ ,  $\beta$ ))
        if v  $\geq \beta$  then return v
         $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
    return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
    if TERMINAL-TEST(state) then return UTILITY(state)
    v  $\leftarrow$   $+\infty$ 
    for each a in ACTIONS(state) do
        v  $\leftarrow$  MIN(v, MAX-VALUE(RESULT(s,a),  $\alpha$ ,  $\beta$ ))
        if v  $\leq \alpha$  then return v
         $\beta \leftarrow \text{MIN}(\beta, v)$ 
    return v
```

The search order is important

- It might be worth while to examine first the successors that are likely to be the best.



Further acceleration?

- Alpha-beta pruning does **not** affect the final result. (good!)
- With perfect ordering, the time complexity is still high, i.e., $O(b^{m/2})$. (bad…)

In practice:

- Replace the **UTILITY** function with a heuristic evaluation function **EVAL**
- Use **CUTOFF-Test** instead of **TERMINAL-Test**

H-MINIMAX(s, d) =

$$\begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = \text{MIN.} \end{cases}$$

III. Generalization of Minimax Algorithm

Stochastics Game



Deterministic game: Scale of the payoff does not matter.



Stochastic game: Scale does matter

Partially Observable Games



Texas Holdem



Starcraft

With partial observability, value of an action depends on the **information state** or belief state the agent is in.

Minimax Optimization

$$\max_{x \in X} \min_{u \in U(x)} f(x, u)$$

- Beyond the tree representation (again).
- An agent can be a general search method.
- Commonly encountered in robust optimization (design, transportation, logistics, etc).

To be continued